



## Same digit interpolates.



## Different digit interpolation



# Code:

```
#Part 1 – Datasets and Dataloaders
```

```
import torch
```

```
!mkdir hw9_data
```

```
from torchvision import datasets, transforms
```

```
root = 'hw9_data'
```

```
transformations = transforms.ToTensor()
```

```
mnist_train = datasets.MNIST(root, train=True, download=True,  
transform=transformations)
```

```
mnist_test = datasets.MNIST(root, train=False, download=True,  
transform=transformations)
```

```
from torch.utils.data import DataLoader
```

```
BATCH_SIZE=32
```

```
kwargs = {'num_workers':1, 'pin_memory':True}
```

```
train_loader = DataLoader(mnist_train, batch_size=BATCH_SIZE, shuffle=True,  
**kwargs)
```

```
test_loader = DataLoader(mnist_test, batch_size=BATCH_SIZE, shuffle=True, **kwargs)
```

#part 2-encoder and decoder:

from torch import nn

from torch.nn import functional as F

inputPixel = 784

outputs = 400

class Encoder(nn.Module):

def \_\_init\_\_(self, latent\_dim):

super(Encoder, self).\_\_init\_\_()

self.fc1 = nn.Linear(inputPixel, outputs)

self.relu = nn.ReLU()

self.fc21 = nn.Linear(outputs, latent\_dim)

self.fc22 = nn.Linear(outputs, latent\_dim)

def forward(self, x):

h1 = self.relu(self.fc1(x))

return self.fc21(h1), self.fc22(h1)

class Decoder(nn.Module):

def \_\_init\_\_(self, latent\_dim):

super(Decoder, self).\_\_init\_\_()

self.relu = nn.ReLU()

self.fc3 = nn.Linear(latent\_dim, outputs)

self.fc4 = nn.Linear(outputs, inputPixel)

self.sigmoid = nn.Sigmoid()

```
def forward(self,x):  
    h3 = self.relu(self.fc3(x))  
    return self.sigmoid(self.fc4(h3))
```

#part 3: TGraining and loss functions:

```
def sample(mu, logvar):
```

```
    std = logvar.mul(0.5).exp_()
```

```
    eps =std.data.new(std.size()).normal_()
```

```
    return eps.mul(std).add_(mu)
```

```
def vae_loss(x, x_hat, mu, logvar):
```

```
    BCE = F.binary_cross_entropy(x, x_hat.view(-1, inputPixel))
```

```
    KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
```

```
    KLD /= BATCH_SIZE * inputPixel
```

```
    return BCE + KLD
```

```
from torch import optim
```

```
encoder = Encoder(latent_dim = 32)
```

```
encoder.cuda()
```

```
decoder = Decoder(latent_dim = 32)
```

```
decoder.cuda()
```

```
params = list(encoder.parameters())+list(decoder.parameters())
```

```
optimizer = optim.Adam(params, lr=1e-3)
```

```
import tensorflow as tf
```

```
def train(encoder, decoder, train_loader, optimizer, num_epochs = 10):
```

```
    train_loss = 0
```

```
    for batch_idx, (data, _) in enumerate(train_loader):
```

```
        data = data.view(data.shape[0],-1)
```

```
        data = data.cuda()
```

```
        optimizer.zero_grad()
```

```
        mu, logvar = encoder.forward(data)
```

```
        z = sample(mu, logvar)
```

```

recon_x = decoder.forward(z)
loss = vae_loss(recon_x, data, mu, logvar)
loss.backward()
train_loss += loss
optimizer.step()
if batch_idx % 10 == 0:
    print('Train Epoch: {} [{}/{}] {:.0f}%]\tLoss: {:.6f}'.format(
        num_epochs, batch_idx * len(data), len(train_loader.dataset),
        100. * batch_idx / len(train_loader),
        loss / len(data)))
    print('====> Epoch: {} Average loss: {:.4f}'.format(
        num_epochs, train_loss / len(train_loader.dataset)))
train(encoder, decoder, train_loader, optimizer, num_epochs = 10)

```

#part 4 Visualizing the VAE output:

```
import matplotlib.pyplot as plt
```

```
from torchvision import utils
```

```
import numpy as np
```

```
def create_interpolates(A, B, encoder, decoder):
```

```
    muA, logvarA = encoder.forward(A)
```

```
    zA = sample(muA, logvarA)
```

```
    recon_a = decoder.forward(zA)
```

```
    muB, logvarB = encoder.forward(B)
```

```
    zB = sample(muB, logvarB)
```

```
    recon_b = decoder.forward(zB)
```

```
    result = [recon_a.detach().cpu().reshape([28,28])]
```

```
    muDiff = (muB-muA)/7
```

```
    muTemp = muA
```

```
    for i in range(7):
```

```
        muTemp += muDiff
```

```
        zTemp = sample(muTemp, logvarA)
```

```
        result.append(decoder.forward(zTemp).detach().cpu().reshape([28,28]))
```

```
    result.append(recon_b.detach().cpu().reshape([28,28]))
```

```
    return result
```

```
similar_pairs = {}
```

```
for _, (x, y) in enumerate(test_loader):
```

```
    for i in range(len(y)):
```

```
        if y[i].item() not in similar_pairs:
```

```
            similar_pairs[y[i].item()] = []
```

```
        if len(similar_pairs[y[i].item()])<2:
```



```

        similar_pairs[y[i].item()].append(x[i])

done = True
for i in range(10):
    if i not in similar_pairs or len(similar_pairs[i])<2:
        done = False

if done:
    break
allP = []
for j in range(10):
    A, B = similar_pairs[j][0], similar_pairs[j][1]
    A = A.view(A.shape[0],-1)
    B = B.view(B.shape[0],-1)
    A = A.cuda()
    B = B.cuda()
    allP.append(create_interpolates(A, B, encoder, decoder))

fig = plt.figure(figsize=(10,9))
for i in range(90):
    fig.add_subplot(10,9,i+1)
    plt.imshow(allP[i//9][i%9])
fig.show()
random_pairs = {}
for _, (x, y) in enumerate(test_loader):
    # Make sure the batch size is greater than 20
    for i in range(10):

```

```

    random_pairs[i] = []
    random_pairs[i].append(x[2*i])
    random_pairs[i].append(x[2*i+1])
    break
# random_pairs[i] contains two images indexed at 0 and 1 that are chosen at random
allP = []
for j in range(10):
    A, B = random_pairs[j][0], random_pairs[j][1]
    A = A.view(A.shape[0],-1)
    B = B.view(B.shape[0],-1)
    A = A.cuda()
    B = B.cuda()
    allP.append(create_interpolates(A, B, encoder, decoder))

fig = plt.figure(figsize=(10,9))
for i in range(90):
    fig.add_subplot(10,9,i+1)
    plt.imshow(allP[i//9][i%9])
fig.show()

```